

Analysis of Variance (ANOVA)

Sumanta Basu

Comparing more than two
means with ANOVA


Outline 1

How about we compare all possible pairs?

- *Pro: Can simply use t-tests many times*
- *Con: Type-I error inflates quickly*
- *Correct for multiple hypothesis testing, at the expense of losing power*

Is there a difference between labs?

Prelims score by lab (not this year)



Should we just compare all pairs?

There are $\binom{4}{2} = 6$ pairs: Could do a two-sample test for each:

- $H_0 : \mu_1 = \mu_2$
- $H_0 : \mu_1 = \mu_3$
- $H_0 : \mu_1 = \mu_4$
- $H_0 : \mu_2 = \mu_3$
- $H_0 : \mu_2 = \mu_4$
- $H_0 : \mu_3 = \mu_4$

Not a good idea.

Multiple testing problem

Suppose we do M independent tests at significance level α . Suppose all nulls are true.

$$P(\text{Fail to reject all } H_0 \mid \text{all } H_0 \text{'s are true}) = (1 - \alpha)^M$$

Suppose we write up a paper on our findings if we rejected any of the nulls. **What's the probability of writing a paper when we shouldn't?**

$$P(\text{Reject at least 1 } H_0 \mid \text{all } H_0 \text{'s are true}) = 1 - (1 - \alpha)^M$$

- E.g. if $\alpha = 0.05$ and $M = 6$, this is 26%.
- if $M = 13$, it's about 50%
- if $M = 30$, it's nearly 80%

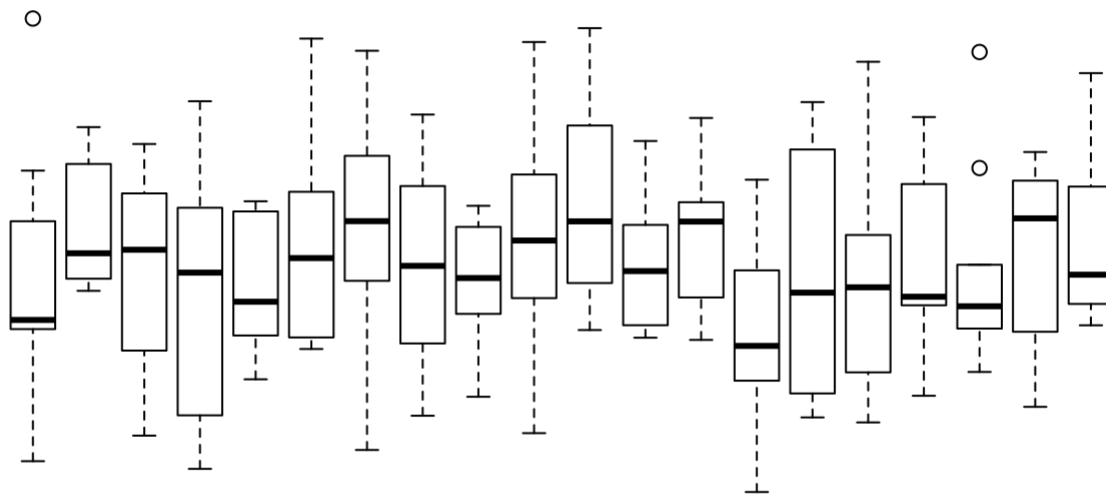
Data snooping / fishing

If you look at the data and then decide which hypothesis to test, your p-value is invalid and you are not actually getting the stated Type I error rate.

It's as if you did some large number of tests M (informally)... so if you chose the "most promising looking" test out of $M = 6$ possibilities, you'd have a 25% of making a Type I error! Even though you used $\alpha = 0.05$.

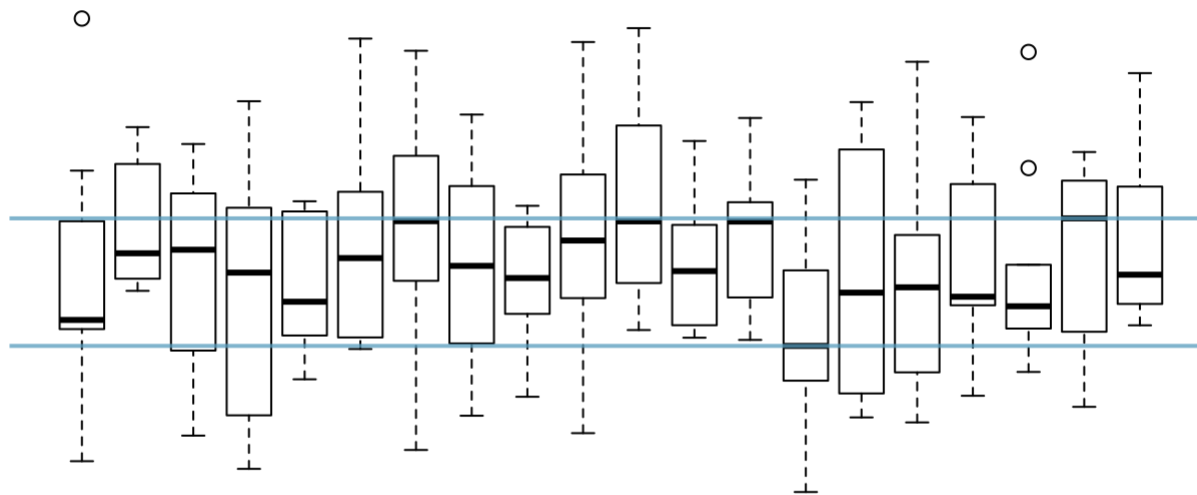
In pictures

Generate 20 groups of size 10 from $N(0, 1)$.



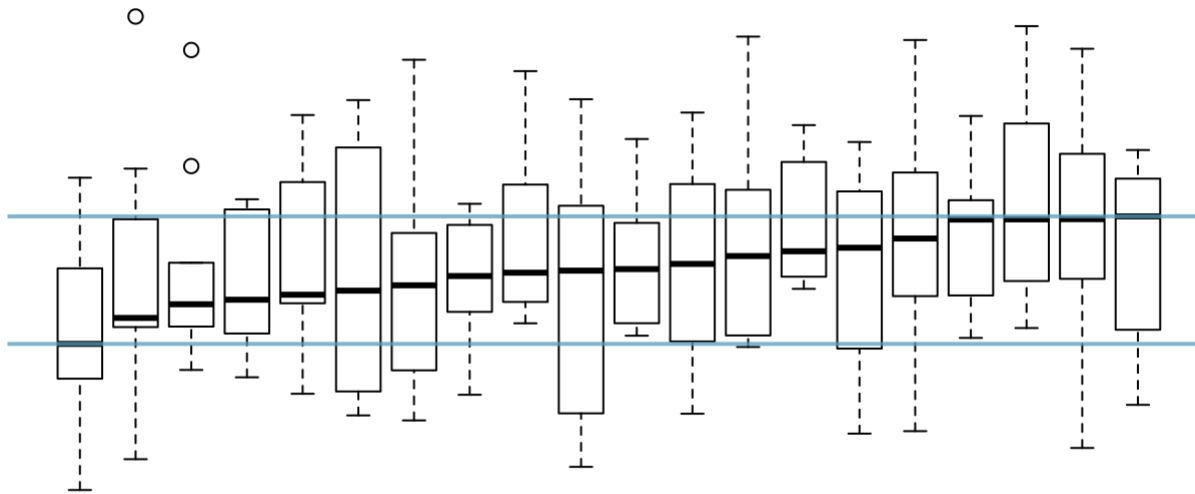
In pictures

A lot of variation in medians... though all from from $N(0, 1)$.



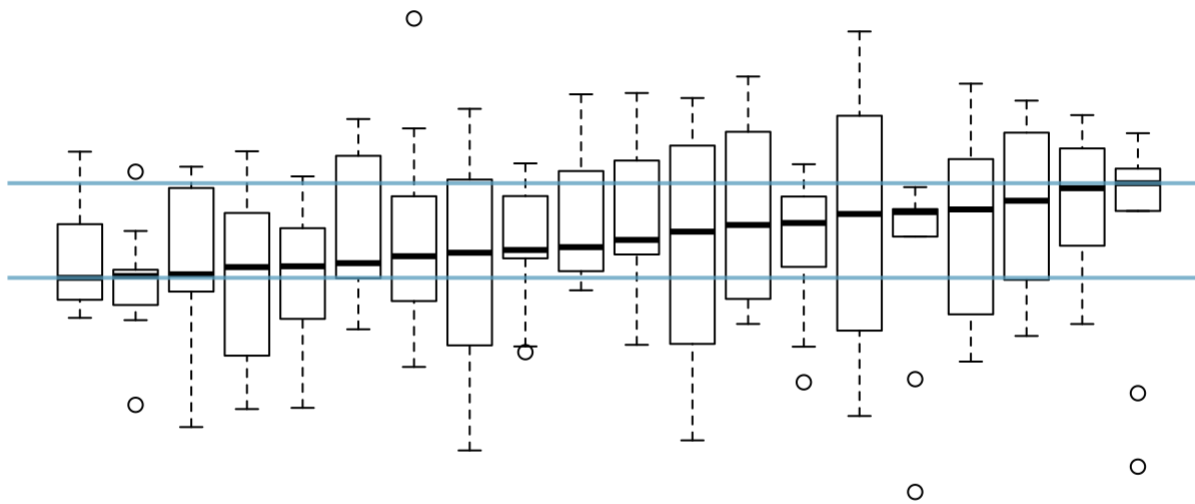
In pictures

Sort these. Lowest dependably far from highest...



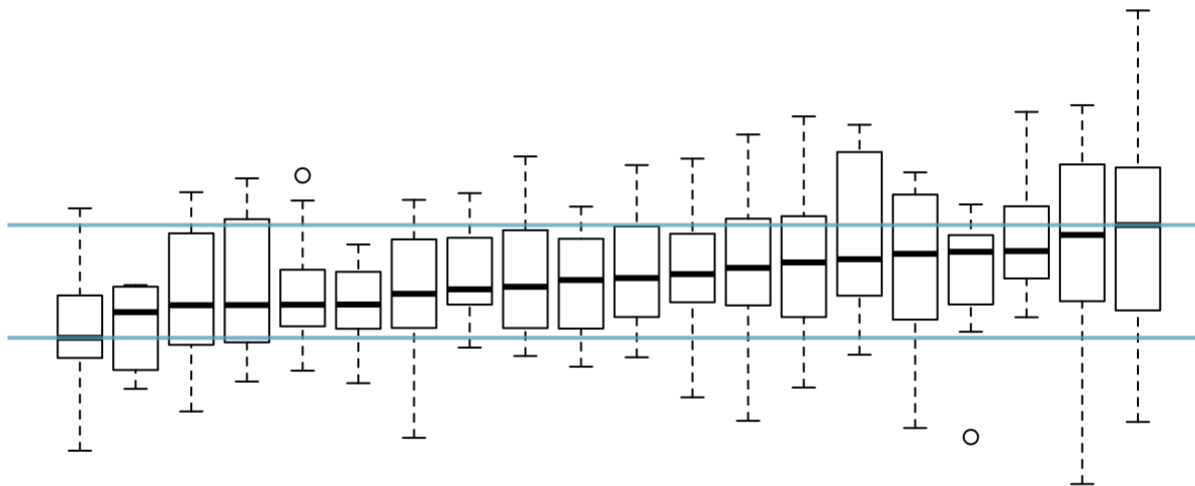
In pictures (repeat)

Sort these. Lowest dependably far from highest...



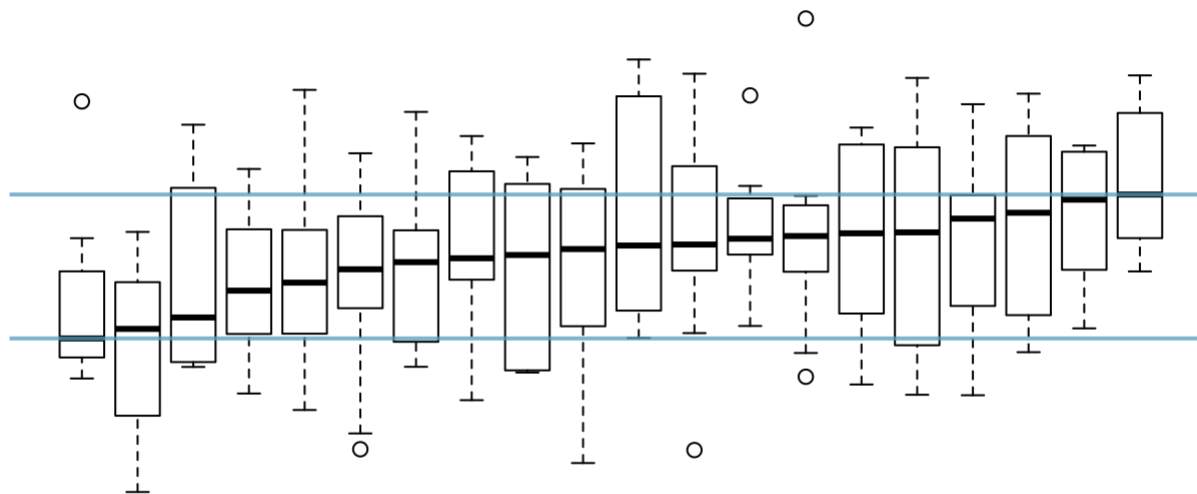
In pictures (repeat)

Sort these. Lowest dependably far from highest...




In pictures (repeat)

Sort these. Lowest dependably far from highest...




In pictures (repeat)

Sort these. Lowest dependably far from highest...



In pictures (repeat)

Sort these. Lowest dependably far from highest...



A simple (but suboptimal) fix

Bonferroni correction If you're doing M tests and you want an overall Type I error rate of α , then do your individual tests at level $\alpha^* = \frac{\alpha}{M}$.

- works even if tests are dependent
- for $M = 6$, we should do tests with $\alpha^* = 0.05/6$.
- not a great fix since hard to reject a null - i.e., we **lose power**
- Bonferroni is simple but overly conservative
- What if you are doing $M = 100,000$ as in genomics??

Test all pairs of labs at 0.05/6 level

Compute p-values: e.g., `t.test(red,blue)$p.value`

```
## [1] 0.9529006
```

```
## [1] 0.4620212
```

```
## [1] 0.9529006
```

```
## [1] 0.3863428
```

```
## [1] 0.9583997
```

```
## [1] 0.4285551
```

- none are even close to $0.05/6 \approx 0.008$

So what should we do?

- after correction for multiple testing, doing all $\binom{k}{2}$ tests can have low power
- sometimes we just want to say that **not all means are equal** without saying which pair it is.
- e.g., lab assignment is *associated* with prelim grades
- not trying to say something specific like lab "Red" is different from lab "Blue"
- We wish to test

$$H_0 : \mu_1 = \mu_2 = \cdots = \mu_k$$

H_A : at least one mean is different

This is the goal of **ANOVA** (**A**nalys**i**s **O**f **V**ariance)

One-way ANOVA

Outline 2

Assume: independence, normality and equal variance

Model each observation as "group mean + noise"

Compare variability between groups to variability within groups

How?

- *Assuming no difference in group means, measure how between/within group variability should look*
- *See if between/within variability is substantially higher in your data*

General guideline: be careful about making causal claims from observational data

Contexts for using ANOVA

Observational study

- draw independent samples of size

$$n_1, \dots, n_k$$

from $k \geq 2$ populations.

- Goal: Are the means of the k populations all the same?
- allows us to make inferences about **associations** (as in correlations)

Controlled experiments

- draw independent samples then randomly assign subjects to one of k treatment groups
- Goal: Are the means of the k populations all the same?
- allows us to make **causal** inferences

Question


Example: Prelim scores in labs

Suppose we reject the null of $\mu_{\text{Green}} = \mu_{\text{Red}}$ in favor of $\mu_{\text{Green}} > \mu_{\text{Red}}$, what can we conclude?
(Confounders?)

0. Assumption and ANOVA model

Assumptions to use ANOVA

- independence (both between groups and within)
- normality
- equal variances



(groups can only differ by mean shifts)

Basic ANOVA: model

Let Y_{ij} = response for unit j from group i (e.g., score of student j in lab i)

All assumptions incorporated by following linear model:

$$\begin{aligned} Y_{1j} &= \mu_1 + \epsilon_{1j} \text{ for } j = 1, \dots, n_1 \\ Y_{2j} &= \mu_2 + \epsilon_{2j} \text{ for } j = 1, \dots, n_2 \\ &\vdots \\ Y_{kj} &= \mu_k + \epsilon_{kj} \text{ for } j = 1, \dots, n_k \end{aligned}$$

where $\epsilon_{ij} \sim N(0, \sigma)$ are all independent.

Response = Mean of group + Noise

Basic ANOVA: model in words

- each group (potentially) has its own mean.
- response is simply the mean corrupted by noise
- noise is independent, $N(0, \sigma)$ with same σ across all groups

Basic ANOVA: model in pictures

Means



Basic ANOVA: model in pictures

Response = Mean + Noise



1. Null and Alternative Hypotheses

ANOVA

$$H_0 : \mu_1 = \cdots = \mu_k$$

What null looks like



Basic ANOVA: model (alternate notation)

$$\begin{aligned} Y_{1j} &= \mu + \alpha_1 + \epsilon_{1j} \text{ for } j = 1, \dots, n_1 \\ Y_{2j} &= \mu + \alpha_2 + \epsilon_{2j} \text{ for } j = 1, \dots, n_2 \\ &\vdots \\ Y_{kj} &= \mu + \alpha_k + \epsilon_{kj} \text{ for } j = 1, \dots, n_k \end{aligned}$$

In this case, we test whether

$$H_0 : \alpha_1 = \dots = \alpha_k = 0$$

Think of α_i as how group i 's mean differs from μ , i.e. $\alpha_i = \mu_i - \mu$

(We require $\sum_{i=1}^k \alpha_i = 0$ so μ represents the "overall mean")

One-way fixed effects ANOVA

$$Y_{ij} = \mu + \alpha_i + \epsilon_{ij} \text{ for } j = 1, \dots, n_i$$

Notation:

- n_i is number of units in group i
- Y_{ij} is response of unit j in group i
- μ is overall mean response value
- α_i is **effect** on mean response due to group i
- ϵ_{ij} random error for unit j of group i

Fixed effects vs. random effects


Fixed effects ANOVA (our focus) - use when we care about the specific k groups present in sample - e.g., we care about these particular four labs

Random effects ANOVA (in BTRY 6020) - use when these specific k groups are not themselves of intrinsic interest, but rather have been chosen to represent a larger population of potential groups - e.g., company wants to know if employee interviews influenced by who's interviewing. SRS of 5 interviewers (out of, say, 100s) and has each interview 10 applicants


2. Measure within and between group variability

Principal idea behind ANOVA


Compare **variability between** the k estimated group-specific means to a pooled estimate of **within-group variability**




Within group variance smaller than
between group variance




Within group variance bigger than between group variance



Within group variance smaller than
between group variance



Within group variance bigger than between group variance



Estimation in fixed effects model

Model: $Y_{ij} = \mu + \alpha_i + \epsilon_{ij}$


Observed data

y_{ij} for $j = 1, \dots, n_i$ (and $i = 1, \dots, k$)

Notation:

- **group-specific sample mean:** $\bar{y}_{i\cdot} = \frac{1}{n_i} \sum_{j=1}^{n_i} y_{ij}$
- **overall sample mean:** $\bar{y}_{..} = \frac{1}{n_{tot}} \sum_{i=1}^k \sum_{j=1}^{n_i} y_{ij}$

(here $n_{tot} = \sum_{i=1}^k n_i$)




Estimation in fixed effects model

Model: $Y_{ij} = \mu + \alpha_i + \epsilon_{ij}$

Obvious estimates:


$$\hat{\mu} = \bar{y}_{..}$$

$$\hat{\alpha}_i = \bar{y}_{i.} - \bar{y}_{..}$$



ANOVA decomposition

$$y_{ij} - \bar{y}_{..} = (y_{ij} - \bar{y}_{i.}) + (\bar{y}_{i.} - \bar{y}_{..})$$



$$\sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..})^2 = \sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2 + \sum_{i=1}^k n_i (\bar{y}_{i.} - \bar{y}_{..})^2$$

ANOVA decomposition

$$\sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..})^2 = \sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2 + \sum_{i=1}^k n_i (\bar{y}_{i.} - \bar{y}_{..})^2$$

$$tss = sse + ssb$$

Total sum of squares = sum of squares within group + sum of squares between groups

- **tss**: variability about the overall mean
- recognize $tss/(n_{tot} - 1)$?

2a. Within group variability

sse (sum of squares of errors)

mse (mean squares of errors)

sse (as pooled error variance)

Regardless of whether H_0 or H_A holds, an unbiased estimate for σ^2 given by

$$\frac{1}{n_{tot} - k} \sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2 = \frac{sse}{n_{tot} - k}$$

Recall **sample variance for group i :**


$$s_i^2 = \frac{1}{n_i - 1} \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2$$

$$"mse" = \frac{sse}{n_{tot} - k} = \frac{1}{n_{tot} - k} \sum_{i=1}^k (n_i - 1) s_i^2$$

mse (as pooled error variance)

$$mse = \frac{sse}{n_{tot} - k} = \frac{1}{n_{tot} - k} \sum_{i=1}^k (n_i - 1)s_i^2$$

- weighted average of groups' sample variances
- bigger groups get more weight



2b. Between group variability

ssb (sum of squares between groups)

msb (mean squares between groups)

ssb: Between groups variability

ANOVA decomposition suggests that

$$ssb = \sum_{i=1}^k n_i (\bar{y}_{i.} - \bar{y}_{..})^2$$

is a way to measure variability **between groups**.

How much *on average* does a group vary from grand mean?

$$msb = \frac{ssb}{k - 1}$$

Under H_0 :

- should be small (purely due to random fluctuations of noise)
- in particular, would expect $msb \approx mse$

3. Compare between to within

- Test Statistic: MSB/MSE
- Under H_0 , follows **F distribution** with appropriate df
- Reject H_0 if test statistic **is large**

F statistic

Under H_0 , we expect $msb \approx mse$. More precisely...

Let MSB and MSE be the random variables (of which msb and mse are realizations)

Under H_0 , the statistic


$$\frac{MSB}{MSE} \sim F_{k-1, n_{tot}-k},$$

that is, it has an F distribution with $k - 1$ and $n_{tot} - k$ degrees of freedom.

F distribution

F stands for Fisher

Density of F with 3 and 120 dfs



Rejecting null

Under H_0 , the statistic

$$\frac{MSB}{MSE} \sim F_{k-1, n_{tot}-k},$$

If $\frac{msb}{mse}$ is "larger than expected under null", we reject H_0 in favor of H_A .

The **ANOVA p-value** is given by

$$P\left(F_{k-1, n_{tot}-k} > \frac{msb}{mse}\right)$$

The ANOVA Table

Source of variability	Degrees of freedom	Sum of squares	Mean square	F statistic	p-value
Between	$k - 1$	ssb	$msb = \frac{ssb}{k-1}$	$\frac{msb}{mse}$	$P(F_{k-1, n_{tot}-k} > \frac{msb}{mse})$
Error (within)	$n_{tot} - k$	sse	$mse = \frac{sse}{n_{tot}-k}$		
Total	$n_{tot} - 1$	tss			

4. Check Assumptions


Independence: revisit sampling strategy used in data collection

Normality: Q-Q plots, histograms





Equal variance: side-by-side boxplots, Levene's Test (to be covered in lab this week)

Back to the labs question

Prelims score by lab




Assumptions met?



A look at Red lab

Histogram of scores[lab == "Red"]



Importance of assumptions for ANOVA

- **independence** is very important
- **normality** is less important especially if sample sizes are close to equal (and not too small)
- **equal σ_i** assumption is crucial
- watch out for **outliers** (can repeat analysis with removed and see if there's a difference)

Let's proceed...

```
fit = aov(scores ~ lab)
summary(fit)
```

##		Df	Sum Sq	Mean Sq	F value	Pr(>F)
##	lab	3	203	67.73	0.359	0.783
##	Residuals	117	22059	188.54		

Rerun after removing outliers

```
fit2 = aov(scores[-ii]~lab[-ii])  
summary(fit2)
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## lab[-ii]      3    334   111.4    0.853  0.468  
## Residuals   112  14631   130.6
```

Still no evidence against null.

5. R commands

aov

```
fit = aov(scores~lab)  
summary(fit)
```

lm

```
scores.lm = lm(scores ~ lab)  
anova(scores.lm)
```

Finishing up ANOVA

Basic ANOVA: model in pictures

Response = Mean + Noise



ANOVA

Want to test

$$H_0 : \mu_1 = \cdots = \mu_k$$

versus H_A : means are not all the same

Basic ANOVA: model

$$\begin{aligned} Y_{1j} &= \mu + \alpha_1 + \epsilon_{1j} \text{ for } j = 1, \dots, n_1 \\ Y_{2j} &= \mu + \alpha_2 + \epsilon_{2j} \text{ for } j = 1, \dots, n_2 \\ &\vdots \\ Y_{kj} &= \mu + \alpha_k + \epsilon_{kj} \text{ for } j = 1, \dots, n_k \end{aligned}$$

In this case, we test whether $H_0 : \alpha_1 = \dots = \alpha_k = 0$

Think of α_i as how group i 's mean differs from μ , i.e. $\alpha_i = \mu_i - \mu$

(We require $\sum_{i=1}^k \alpha_i = 0$ so μ represents the "overall mean")

One-way fixed effects ANOVA


$$Y_{ij} = \mu + \alpha_i + \epsilon_{ij} \text{ for } j = 1, \dots, n_i$$

Notation:

- n_i is number of units in group i
- Y_{ij} is response of unit j in group i
- μ is overall mean response value
- α_i is **effect** on mean response due to group i
- $\epsilon_{ij} \sim N(0, \sigma)$ random error for unit j of group i

Back to the labs question

Prelims score by lab



In R...

```
scores.lm = lm(scores ~ lab)
anova(scores.lm)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: scores
```

```
##           Df  Sum Sq Mean Sq F value Pr(>F)
```

```
## lab         3    203.2   67.729   0.3592 0.7826
```

```
## Residuals 117 22059.0  188.539
```

Rerun after removing outliers

```
scores2.lm = lm(scores[-ii] ~ lab[-ii])  
anova(scores2.lm)
```

```
## Analysis of Variance Table  
##  
## Response: scores[-ii]  
##           Df  Sum Sq Mean Sq F value Pr(>F)  
## lab[-ii]    3   334.1   111.37   0.8526 0.4681  
## Residuals 112 14630.6   130.63
```


Still no evidence against null.

6. If H_0 rejected, which pairs differ in mean?

- Can go back to pairwise t-tests, with multiple testing correction (using, say, Bonferroni)
- Can use Tukey's Honest Significance Difference (HSD) test

Suppose the data looked like this

(Fake) prelims score by lab



ANOVA on fake data

```
scores3.lm = lm(scores~lab, data = dat2)
anova(scores3.lm)
```

```
## Analysis of Variance Table
##
## Response: scores
##           Df  Sum Sq Mean Sq F value  Pr(>F)
## lab         3  1121.3   373.77   2.8613 0.04007 *
## Residuals 112 14630.6   130.63
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- At $\alpha = 0.05$ level, we reject H_0 in favor of H_A .

Which pair(s) are different?

- can do pairwise tests as long as we correct for multiple testing
- in this context, there's something better than Bonferroni...
- **Tukey Honest Significant Differences** (Tukey HSD)
- takes into account that we are looking at $\binom{k}{2}$ comparisons

In R

```
library(multcomp)
hsd <- glht(scores3.lm, linfct = mcp(lab = "Tukey"))
summary(hsd)
```

[Output on next slide.]

In R

```
##
## Simultaneous Tests for General Linear Hypotheses
##
## Multiple Comparisons of Means: Tukey Contrasts
##
##
## Fit: lm(formula = scores ~ lab, data = dat2)
##
## Linear Hypotheses:
##              Estimate Std. Error t value Pr(>|t|)
## Yellow - Green == 0    3.294      3.208   1.027   0.7336
## Red - Green == 0       4.325      2.927   1.477   0.4537
## Blue - Green == 0      8.370      2.883   2.903   0.0225 *
## Red - Yellow == 0      1.031      3.186   0.324   0.9882
## Blue - Yellow == 0     5.076      3.146   1.613   0.3746
## Blue - Red == 0        4.045      2.859   1.415   0.4920
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## (Adjusted p values reported -- single-step method)
```

Final note

- can sometimes reject $H_0 : \mu_1 = \mu_2 = \dots = \mu_k$ and yet no pairwise test would reject!
- ANOVA pools information across all groups and tries to answer a more general question