Data collection principles

Sumanta Basu

Reading

- Reading: Textbook Sections 1.1, 1.2.1, 1.3 - 1.5
- Recommended Exercise: 1.1, 1.9, 1.13, 1.17, 1.27, 1.29
- OpenIntro Lab 0: Introduction to R (under blackboard folder ‘Lab1’)

Data Sampling

main focus of course will be on

- data analysis
- statistical inference
- careful interpretation of results

but **how you collect data** is essential (e.g., Agent Orange example)

many contexts for data collection

- surveys
- observational studies
- scientific experiments

Sampling: a crucial statistical concept

historically, statistics concerned gathering census-style data on entire population of interest (“state”)

> “You don’t have to eat the whole ox to know the meat is tough.” – Samuel Johnson

why not examine entire population?

- expensive, slow
- may be infeasible (e.g., if observation destroys what’s being measured)

modern notion of statistics: learn about the population by examining a small fraction of it

Statistical Sampling

Terminology

Population

- the collection of all units of interest
- implied by the research question being studied
- e.g., all Cornell graduate students

Sample

- any subset of the units from the population
- e.g., 20 Cornell graduate students

Unit or **subject** or **element**

- individual entity from the population
- e.g., a Cornell graduate student

Example

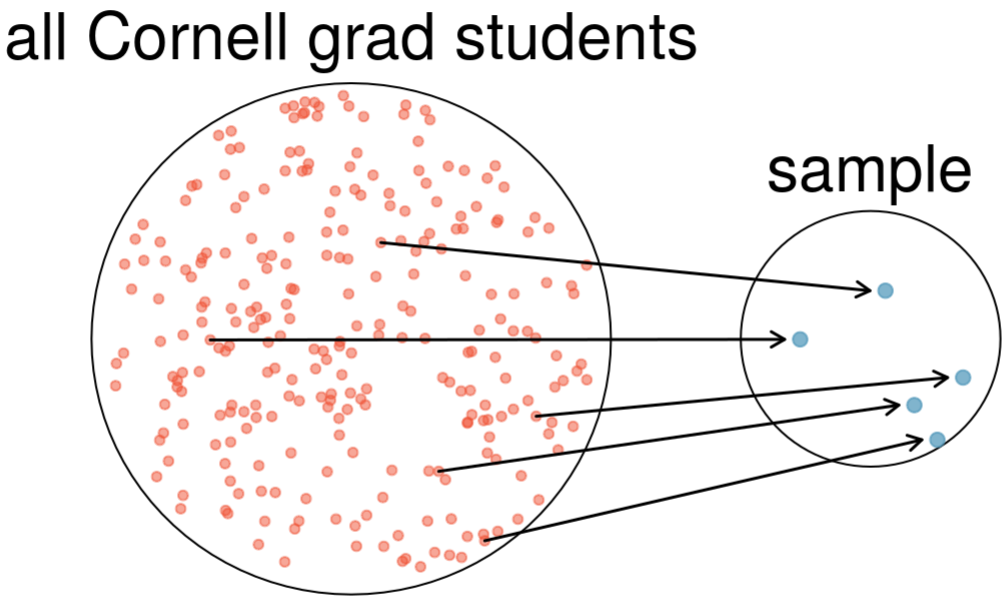

Some distinctions

Randomness

- Population is **fixed / non-random / deterministic**
- Samples could be (often are) **random**

Observability

- Population can be (often is) unobservable
- Sample is observable (can measure and do analysis with it)

What is the **source of randomness** in a sample?

- Sampling Scheme / Sampling Strategy !!

How do we choose the sample?

- Critically important question
- A bad sampling strategy can invalidate all conclusions
- A good sampling method ensures that samples are (with high probability) **representative of the population**
- Choice of sampling method depends on objective of study and feasibility

Sampling design process

Define target population

- What collection of units would you like to describe?
- Implied by the research question being studied
- Population exclusions? (e.g., patients too ill for study)

Example

Research question: What proportion of Cornell graduate students entered graduate studies immediately after graduating college?

- Population?
- The unit?
- The sample?
- Exclusions?

Determine sampling frame

enumerate the population - i.e., a “list” of the population which will help you reach the sample

sampling units may be households, streets, telephone numbers, fields, etc.

example: if population is all Cornell students, could use

- registrar
- university directory

common problems with lists: access, omissions, out-of-date, duplicates

Selecting a sampling design

Probability sampling

- all units in population have a known (non-zero) probability of being included in sample
- simple random sampling
- stratified sampling
- cluster sampling

Non-probability sampling

- some units have zero chance of being included or are included with unknown probability
- convenience sampling, judgment sampling, snowball sampling, quota sampling

Probability sampling

Probability sampling

- aka **random sampling**
- objective procedure
- **probability** of selecting each unit is **nonzero** and **known in advance**
- **ensures representative** sample of the population
- sampling error can be assessed
- results can be projected to population
- more expensive than non-probability samples

Simple random sampling (SRS)

- like “picking names out of a hat”
- an SRS of size n is such that **all possible samples of size n are equally likely**

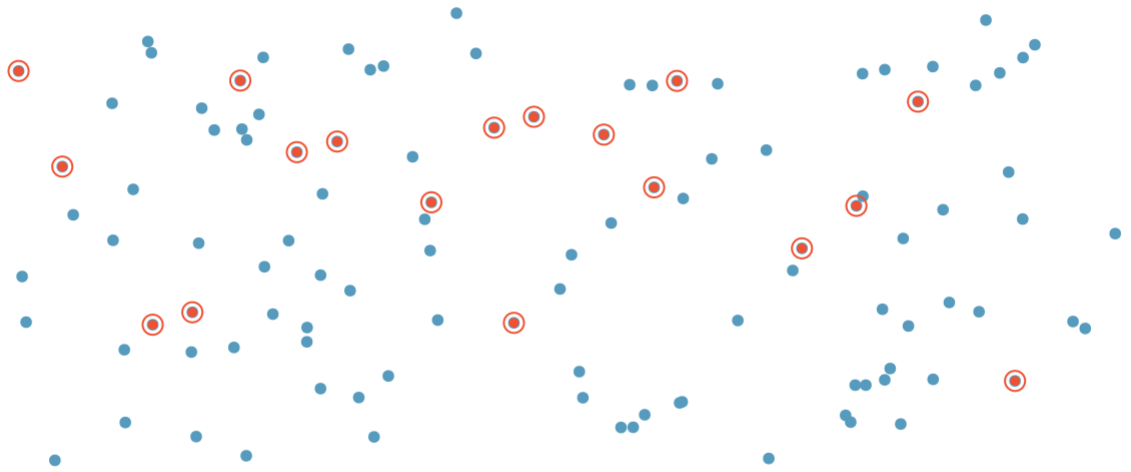

Is this an SRS of 6010 students?

Include in sample if last name starts with the letter “M”.

SRS of 6010 students

- number students from 1 to 150
- now, choose 10 at random (without replacement)

```
sample(1:150, size=10, replace=FALSE)
```

```
##  [1]  88 134  86  47  71  87  10 106  92 144
```

```
sample(1:150, size=10, replace=FALSE)
```

```
##  [1] 105  53  47 145 110 103  33  37  93  88
```

SRS

- simplest method that gives valid inferences
- intuitive, yet sophisticated, idea
- simplest calculations
- basis of most good sampling methods
- workhorse of the sampling world

SRS

- but is it safe to assume that SRS is representative of whole population?

> Theorem: When a sample is selected by SRS, it differs very little from the entire population.

(We'll make this precise later.)

- This course primarily focused on inference based on data obtained from SRS

Stratified random sampling

chosen sample forced to contain units from each “statum” of population

- goal: to equalize “important” variables (e.g., gender, race, geographical area, etc.)
- procedure:
  - partition population (strata are mutually exclusive and exhaustive) based on a characteristic
  - draw simple random samples within each stratum

Stratified random sampling

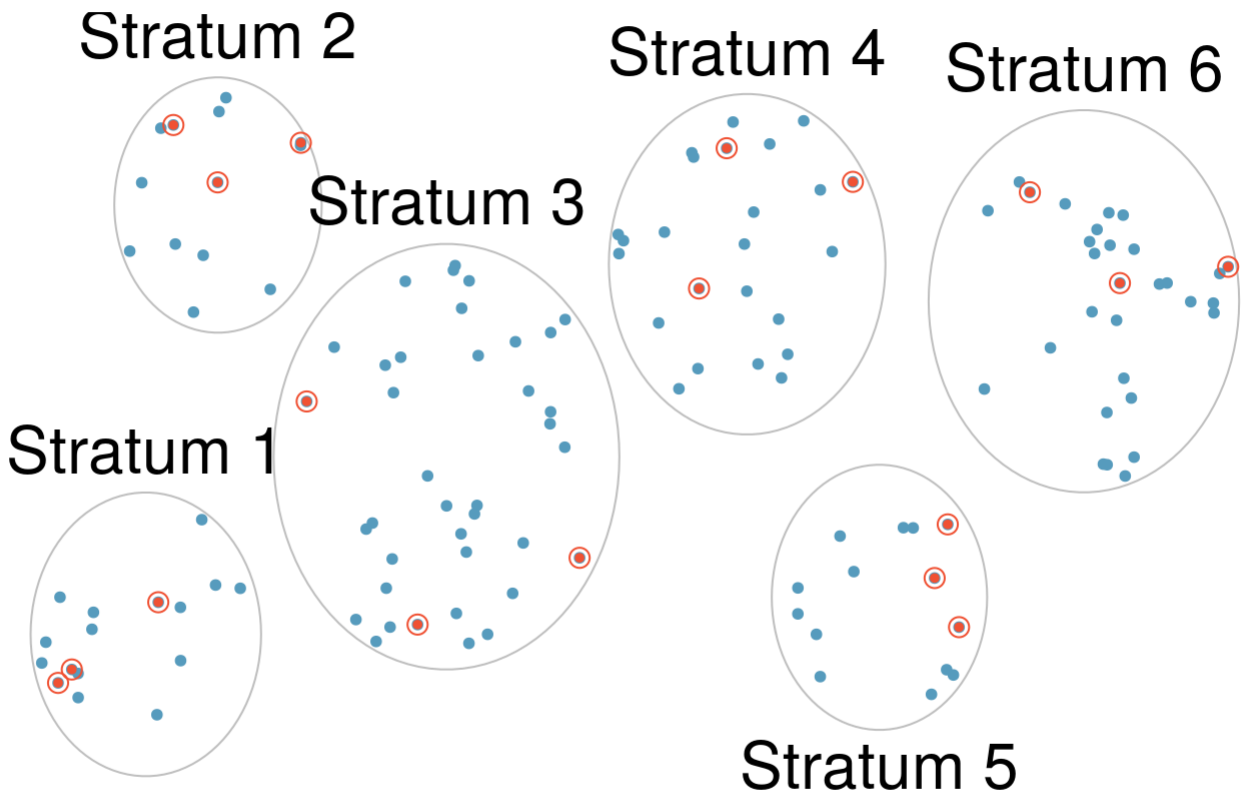

Stratified random sampling

- smaller sampling error than SRS (removes a source of variation)
- representativeness when **proportional sampling** used
- e.g., 6010 students:

Proportional Sampling

	Social Science	Science
Population	50	100
Sample	5	10

- **disproportional stratified sampling** used when some strata are too small OR more important OR more diverse

Cluster sampling

- partition population into clusters (ideally, each adequately represents population)
- SRS of a few clusters
- within each chosen cluster do SRS

Motivation: more economical

Cluster sampling

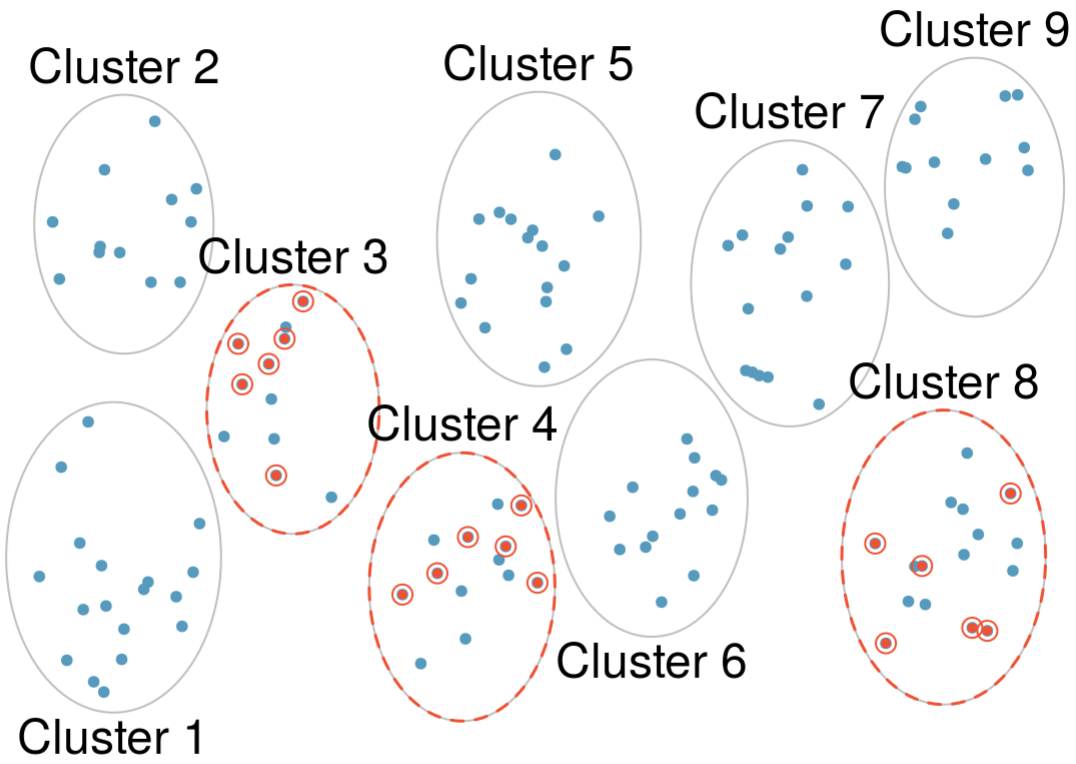

Example

- cluster sampling **by geographic region**
- divide up map, sample regions
- within each region sample units
- can be more economical

Selecting a sampling design

use **stratified** sampling

- when primary objective is to compare groups
- when it may reduce sampling error (homogeneous strata)

use **cluster** sampling

- when there are fixed costs associated with each data collection location
- when list of clusters is available but no list of units themselves

Non-probability sampling

Non-probability sampling

- subjective procedure in which **probability of selection for some units is zero or unknown** before drawing the sample
- sampling error cannot be computed
- non-representative sample so results *cannot* be projected to the population
- cheaper and faster, but allows for limited inference

Convenience sampling

- send survey to your friends on facebook
- stop people on the street
- student volunteers (e.g., undergrad psych students!)

Example

Example

Slate (2013):

> “Sixty-seven percent of American psychology studies use college students, for example. This means that many or even most of the subjects are teenagers.”

Judgment sampling

- researcher tries to select a sample that seems to be most appropriate for study
- e.g., sampling from a particular shopping mall that researchers “think” is representative of their target market

Snowball sampling

- initial subjects recruit additional ones (“referral sampling”)
- used to sample difficult to reach populations (e.g., drug users, sex workers, etc.)
- biased (e.g., over-represents people who have a lot of friends)

Quota sampling

Procedure:

- partition population based on a characteristic (as in stratified random sampling)
- establish how many from each group must be sampled
- non-randomly choose required number from each group

Non-probability sampling

- may be less time consuming or less expensive
- if certain elements have zero probability, then sampling does not fully represent population
- if probabilities of being in sample are not known, cannot project estimates from sample to population
- drawing inferences with a non-representative sample is dangerous (e.g., TV viewer polls)

TV viewer polls

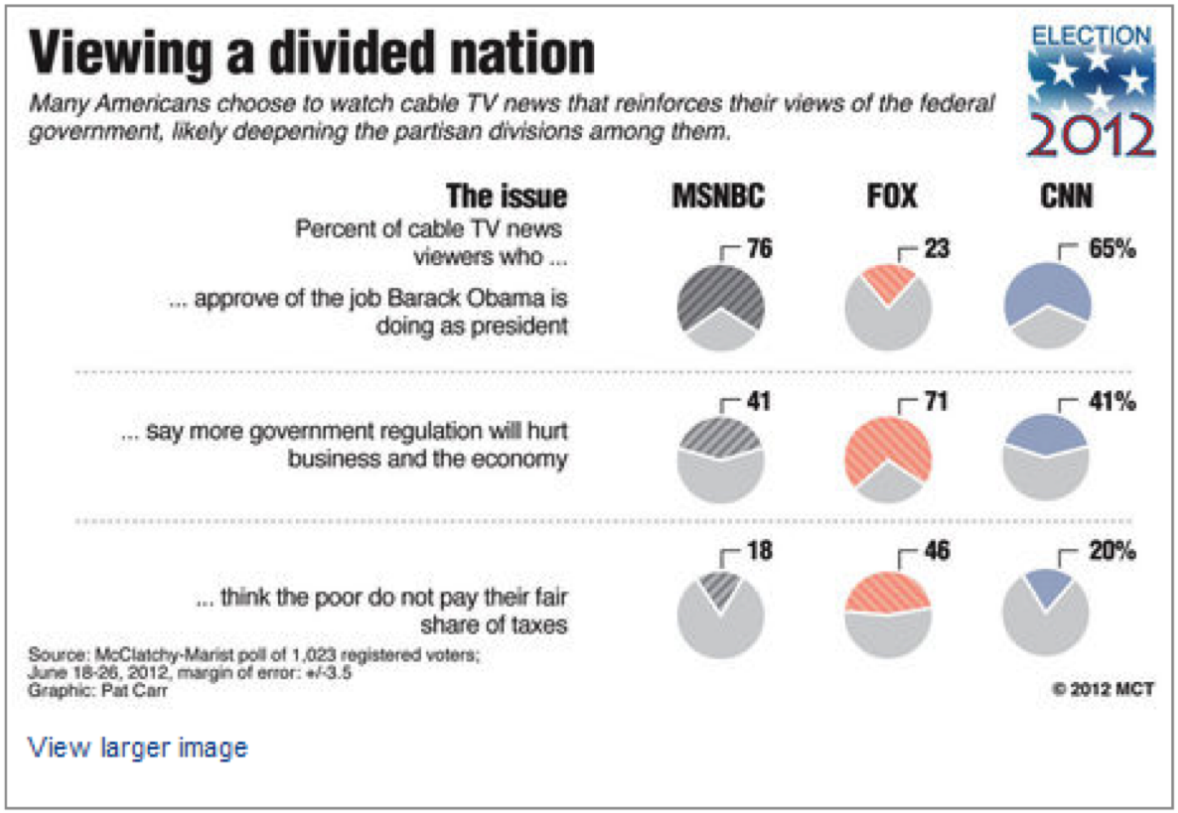

http://www.mcclatchydc.com/2012/07/26/157899/fox-mnsbc-viewers-see-world-differently.html

Errors in sampling

Random sampling error

- sample selected is not representative of population due to chance
- controlled by sample size (larger sample means less likely to have non-representative sample)

Non-sampling error

- systematic error
- **not controlled by sample size**

Types of non-sampling errors

Non-response error

- problematic if non-responders and responders are different

Response or data error

- systematic bias during data collection, analysis, or interpretation
- respondent error (lying, forgetting, etc.)
- interviewer bias
- recording errors
- poorly designed questionnaires

Example: 1948 presidential election

Three major polls predicted Dewey would win.

They were wrong…

What went wrong?

- stopped polling too soon (two weeks before election)
- telephone polls oversampled wealthy
- in person polls used stratified sampling to obtain samples representative of overall US population in terms of race/gender. Not representative of US **voting** population.
- Pollsters used quota sampling

Example: Hormone replacement therapy

1991: Every woman should get **on** HRT immediately because it prevents heart attacks! (50% reduction)

2002: Every woman should get **off** HRT immediately because it increases risk of heart attacks! (30% increase)

What happened?

- 1991 data from the *Nurses' Healthy Study*, which could only answer the question "What are the health characteristics of nurses who choose to undergo HRT therapy versus nurses who don't?"
- 2002 data came from randomized clinical trial
- **selection bias** (healthy volunteer problem)

Experiments

- in the best case, one can randomly select which units are assigned a treatment
- one then compares units in the treatment group to those in the control group

Four fundamental principles in designing experiments

1. **Controlling**: attempting to ensure that units are treated identically except for treatment
2. **Randomization**: randomly assigning units to treatment/control "evens out" any uncontrollable differences and prevents accidental bias
3. **Replication**: increased sample size reduces statistical fluctuations
4. **Blocking**: stratify based on a likely-relevant variable (e.g., gender) and then randomly assign treatment/control within each block.